ALTERNATIVE IMPLEMENTATIONS OF LARGE SCALE GIS PROCESSING

by

Nathan Thomas Kerr

A Thesis Presented in Partial Fulfillment
of the Requirements for the Degree
Master of Science

ARIZONA STATE UNIVERSITY

August 2009

ALTERNATIVE IMPLEMENTATIONS OF LARGE SCALE GIS PROCESSING

by

Nathan Thomas Kerr

has been approved

August 2009

Graduate Supervisory Committee:

Dr. Daniel Stanzione, Chair
Dr. Robert Pahle
Dr. Yi Chen

ACCEPTED BY THE GRADUATE COLLEGE

ABSTRACT

Abstract goes here

Dedication goes here

ACKNOWLEDGEMENTS

Acknowledgments go here.

TABLE OF CONTENTS

LIST OF TABLES

LIST OF FIGURES

CHAPTER 1   INTRODUCTION

Geographic information systems (GIS) were designed to describe aspects of the world around us. GIS data is made up of geometries such as points or a polygons at certain longitudes, latitudes, and altitudes along with related descriptive information such as a land use type. GIS data can be used to represent a large range of real-world objects such as road or power networks, building locations, or natural features such as lakes and rivers. Utilizing GIS data is one method toward processing, analyzing, and simulating real-world systems. One consumer application of GIS are the GPS based car navigation systems common today.

Larger scale GIS applications also exist in areas such as city planning. City planners use GIS to study road networks, zoning issues, and to simulate population growth. As cities and metropolises grow, the amount of data required to represent these areas also increases. As the amount of data increases, so does the processing power required to complete the geospatial analysis, processing, and simulation.

Geospatial (GIS) simulation and analysis are important processes to understanding and improving our environment, both urban and natural.

1.1   Current Methods

Desktop GIS packages such as ArcGIS[1], QuantumGIS[2], and GRASS GIS[3] are commonly used for GIS processing and analysis. While these programs provide graphical interfaces to their GIS capabilities, their capabilities are limited by the computers they run on. Datasets can be too large for their memories and computations can take too long to be practical. The popular urban simulation package, UrbanSim[4] faces these same constraints.

An alternative approach to using these desktop programs is to employ a geospatial

database like PostGIS[5], ArcSDE[6] or Oracle Spatial[7]. Geospatial databases allow centralized access to, and processing of, geospatial data through query languages such as SQL. As data is stored and managed by the database software, advanced database features such as indexes can be utilized to speedup data access and processing.

PostGIS is utilized as the core component of the Urban Systems Frame-work[8] (USF) designed by the Digital Phoenix[9] project group at Arizona State University. Digital Phoenix tries to integrate 3D visualization technology with simulated and gathered GIS data to better understand the impacts of urban planning decisions.

1.2 Limitations

GIS data has become easier to gather in recent years with the proliferation on low cost GPS devices. The availability of these devices significantly reduces the cost of data collection, moving it from a government funded service to the capabilities of private companies and individuals even to the point of open source style maps such as OpenStreetMap[10].

With the reduced cost of data collection, the amount of data has grown significantly. As datasets grow and the associated processing becomes more sophisticated, the limitations of programs that only work on a single computer are felt through long processing times and inability to work with the required data. Thus a method of utilizing multiple computers to complete the required processing is needed to increase the size of the data that can be processed and reduce the time required to complete the processing.

1.3 Attempted solutions

Attempts have been made to overcome the limitations of single machine implementations. Of primary interest are those extending current methods to use multiple computers.

1.3.1 Parallel Databases

One method of using multiple computers to perform the required processing is the use of parallel databases[11]. Parallel databases should be able to spread both data storage and processing across multiple computers transparently from the view of the SQL query programmer. Parallel database techniques have been used to build a scalable geospatial database system[12, 13].

Data is spread between computers using round-robin, hash, or spatial partitioning. Because the data is able to be distributed between multiple computers, the processing is able to scale to larger datasets. When a query is processed across the database, a thread is created for each fragment of the data. Thus as the data grows larger, the processing capabilities of the system also increase.

Databases excel at working with indexed data while allowing multiple users to interact with the data in a concurrently safe manner through the use of atomic transactions. The requirements placed upon database systems to handle these situations slow down computations that don't utilize indexes or work on an entire dataset at once. The processing operations this research examines do not require these restrictions, and as such a more efficient system can be created.

1.4 Problems with Desktop Programs on Clusters

Current desktop approaches to GIS processing such as ArcGIS are unable to make use of the multi-machine processing environment that compute clusters provide. While reworking these programs to utilize these extended resources is possible, it is non-trivial. GRASS GIS was reworked[14] to use its collaboration features to distribute sub-queries among computers. The method described in this paper utilizes multiple instances of GRASS in a master-slave configuration where all participants access a shared data repository or

filesystem. The geometries are portioned between the various nodes. Operations are done on the subsets, and the results are merged to produce the final result.

While the method used to extend GRASS GIS will in fact speed up GIS processing, it has two flaws. First, GRASS is designed to be used in an interactive mode rather than a batch or script driven approach. Second, the entire set of data used must still fit on a single computer to move it in or out of the environment.

1.5   Cluster Programming Model

Many universities and research institutions already have significant investment in compute clusters. These clusters are groups of computer linked together with high speed network interconnects and high performance parallel filesystems such as Lustre[15].

Cluster access is a valuable resource, and is treated as such. To maximize usage, cluster resources are allocated through batch queuing systems. In general, interactive access is not allowed. Therefore, programs that run on clusters need to be able to work without user input at the time of execution.

Programs that run on clusters require a method of utilizing the computers allocated. A common method in use today is message passing using MPI. MPI is a standard with various different implementations, one of which is installed on most clusters. Moving from one MPI implementation to another usually just requires recompiling the program.

Cluster based programs can be evaluated in terms of speedup and scalability. Speedup is the amount of extra processing power gained by utilizing multiple processors. Scalability describes the program's ability to operation on various numbers of processors utilizing various amounts of data.

1.6   Cluster based approaches to GIS

This thesis implements and evaluates two parallel, dataset centric approaches to processing large geospatial datasets on clusters. The first approach, called HadoopGIS, uses the Hadoop[16] map/reduce framework. The second approach, ClusterGIS, uses the more traditional approach to programming for clusters, MPI. Both approaches provide the geospatial operations required by the Open Geospatial Consortium's Simple Features[17] standard. By applying data parallel programming methods and dispensing with record centric processing methods, both these methods create a fairly easy environment to program in while providing significant speedup and scaleup.

This document first discusses related works in GIS processing and paralllization. These related works are used to form the requirements of a good cluster based implementation. These requirements are then transferred into the design of HadoopGIS and ClusterGIS. After the implementation design discussion is concluded, the evaluation method is discussed in the experimental setup. From the results of the experiments, the performance of the implementations are evaluated and analyzed in the results section. Then I provide some recommendations and conclusions.

CHAPTER 2   RELATED WORK

GIS processing is not new, and a large body of work exists to drawn from. Here I discuss a few aspects of that body of work that are most applicable to this thesis. I start with the programmer's view of GIS - libraries. Then I move to serial and parallel shared-memory programs. Then I move to work done on parallel, distributed memory implementations. The main goal of parallel applications is to make use of multiple processors to complete the required processing more quickly. In short, the processors must work together to complete a single processing operation.

2.1   GIS libraries

The Open Geospatial Consortium (OGC) defined a core set of geospatial processing operations in their Simple Features[17] standard. These core operations allow for most geospatial processing needs. One of the main libraries that provides these operations and related data types is the Java Topology Suite[?] (JTS). Though written in Java, the JTS has been used as the basis for ports into other languages. The Geometry Engine, Open Source (GEOS) library is a C++ port of the JTS that also provides a C interface. PostGIS is implemented using the GEOS library. The NetTopologySuite[?] (NTS) is a port of the JTS into .NET.

As quality libraries are available for a variety of languages, there is no need to reimplement the functionality they provide. This thesis makes direct use of the JTS and GEOS libraries.

Figure 2.1: Shared-Memory Machine Architecture

2.2 Serial and Parallel Shared-Memory Applications

Computer programs are executed by processors. The simplest of programs is made up of a series of operations that are executed in order by the processor. As this type of application is comprised by a series of operations it is known as a serial program. Serial programs are not able to take advantage of more than one processor. To utilize more than one processor at a time, a set of serial programs that can communicate with each other is required. Each serial program in this set is referred to as a thread. Thus multi-threaded programs can utilize more than one processing core. The simplest method of communication between threads is to share a common section of memory. Therefore a parallel shared-memory application could utilize at most the number of processors able to be connected to a single section of memory.

Both serial and parallel shared-memory applications are limited to a single computer, where computer means a group of processors that are connected to the same memory. Most modern computers provide this capability.

The wide variety of desktop GIS processing applications such as ESRI's ArcGIS[1], QuantumGIS[2] and GRASS GIS[3] work on shared memory machines. While these programs provide GIS processing capabilities and graphical user interfaces, they are limited to working on a shared memory. Thus they are limited to a single computer's memory and processing capabilities.

Geospatial databases such as PostGIS[5] and ArcSDE[6] also work on shared-memory machines, but with the intent of sharing the processing capabilities of the machine they run on. Because of my work on the Urban Systems Framework[8] (USF), I am most familiar with PostGIS. PostGIS is able to execute multiple queries at one time, but no

Figure 2.2: Distributed-Memory Machine Architecture

single query uses more than one thread. Even if these databases could utilize all the processing capabilities of the computer to execute one query, they would still be limited to a single machine.

2.3 Parallel Distributed-Memory Applications

To overcome the limitations of a single computer, the more scalable architecture of a parallel distributed-memory machine was created. The basic unit of this architecture is a processing element (PE) comprised of one or more processors couple with memory. In other words, a PE is a shared-memory machine. The PEs are interconnected using some sort of networking technology. This architecture scales as well as the interconnect does. Common interconnect technologies in use today are Gigabit-Ethernet and InfiniBand. Clusters are parallel distributed-memory machines.

For a program to run on a parallel distributed-memory machine, it must be able to work with multiple thread of operation where communication between the threads goes over the interconnect. For the purposes of discussion, threads running on different PEs are called tasks.

The main task in designing a parallel program is figuring out how to split up the work between tasks. One method is to split up the processing between tasks. Each task would generally have the same data and perform variations of an operation on the data; for instance, running multiple scenarios. This methodology is called Task Parallel. An example of task parallelism is simulating the effects of various weather patterns on an urban environment. Each task would have a copy of the environment and run its variety of weather on it. The capability of executing each scenario is limited to the capabilities of

a single processing element, but many scenarios can be executed at the same time.

Data Parallel methods split the data up between tasks an perform the same, or similar, operations on each piece of data. To calculate the effects of a weather pattern on an urban environment, the environment would first be divided between the tasks with each task responsible for one part of the environment. Each task would then calculate the effects of the weather on its section of the environment, communicating with the other PEs as needed to share information related to edge conditions, etc.

2.3.1   Parallel Databases

Parallel databases such as TeraData[18] and Oracle[7] use data parallel methods. Paradise[12, 13] spreads data between computers using round-robin, hash, or spatial partitioning. Because the data is able to be distributed between multiple computers, the processing is able to scale to larger datasets.

When a query is processed across the database, a task is created for each fragment of the data. Thus as the data grows larger, the processing capabilities of the system also increase. If a particular processing operation requires relatively less computation for each data record this is fine. However, after the amount of computation per data record increases beyond a certain point, which is dependent on the speed of the machine used, the processing operation can be sped up by utilizing more processors. Major factors in determining how much data should be processed on each machine, and therefore how the data should be spread between machines, are memory, computation, and communication overhead to move the data and computation to another machine. The ratio of computation to memory and commutation requirements is often referred to as grain size. Coarser grained processes have more computation per data record, while finer grained processes have little computation for each data record.

Databases excel at working with indexed data while allowing multiple users to interact with the data in a concurrently safe manner through the use of atomic transactions. The requirements placed upon database systems to handle these situations slow down computations that don't utilize indexes or work on an entire dataset at once. The processing operations this research examines do not require these restrictions, and as such a more efficient system can be created.

Many universities and research institutions already have significant investment in compute clusters. These clusters are groups of computer linked together with high speed network interconnects and high performance parallel filesystems such as Lustre[15]. By separating compute and storage resources at the cost of a high speed network, compute clusters are able to separate computation from data storage.

Parallel filesystems allow the data to be separated from the computer where the processing will be executed by spreading files across multiple network connected fileservers allowing access that can be faster than utilizing a computer's local disk for storage while also enabling processing to spread across the available compute resources based entirely on the process' grain size.

GRASS GIS was reworked[14] to use its collaboration features to distribute subqueries among computers. The method described in this paper utilizes multiple instances of GRASS in a master-slave configuration where all participants access a shared data repository or filesystem. The geometries are portioned between the various nodes. Operations are done on the subsets, and the results are merged to produce the final result.

While the method used to extend GRASS GIS will in fact speed up GIS processing, it has two flaws. First, GRASS is designed to be used in an interactive mode rather than a batch or script driven approach. Second, the entire set of data used must still fit on a single computer to move it in or out of the environment.

batch mode processing

data distribution options

standard geospatial operations

scalable, high performance

dataset centric

A good cluster based GIS processing environment needs to be separate from a graphical user interface so that it does not incur additional overhead on the compute nodes and so that it can interact well with the generally batch-scheduled environment of a cluster. In addition, it must be able to distribute data between all the nodes used such that a single node does not become a limiting factor in the environment's scalability. A variety of data distribution models should also be available so that the data is distributed in a manner consistent with the required processing. Of course, the environment should be able to execute all the standard geospatial operations as defined by the Open Geospatial Consortium in their Simple Features[17] standard.

3.1   easy to use to form new operations

Sample program:

initClusterGIS

loadData (distributed/replicated)

process (user code here)

saveData (distributed/replicated)

finalizeClusterGIS

CHAPTER 4   DESIGN

How to fulfill requirements

Architecture of solution

Split dataset across tasks (MPI I/O)

Each task works on their own part

MPI to communicate if needed

CHAPTER 5   EXPERIMENTAL SETUP

How to show how well requirements were filled

Some of the needed research has already been completed. After determining the query set and creating the PostGIS implementation of it, the processing environment was prototyped in Hadoop[16], an opensource implementation of Google's MapReduce[19] framework. MapReduce works by utilizing two user defined function, map and reduce. For each record of the primary dataset the map function is called. This process is made parallel by splitting the data up among several computers and running the individual map functions on those computers. After the map function is complete, the resulting set of data is passed to several reduce functions. Each instance of the reduce function receives all the data from the resulting dataset associated with the same key. The lessons learned from this prototyping phase will be applied to the MPI implementation.

5.1   Operation set

These operations are representative of problem space.

Data access based: parallelization is based off data decomp, so operations are too.

OGC standard requirements - survey of methods requiring 1, 2, or more data points.

http://www.opengeospatial.org/standards/sfa (Simple Feature Access - Part 1: Common Architecture 1.2.0)

All operations must be performed as if they were part of a series of operations.

1. Create a new record

PostGIS: insert... Hadoop: map: emit all record giving them the same key (forcing one reduce); reduce emits all received records with the addition of one ClusterGIS: emits

Section	1 Geometry	2 Geometries
6.1.2 (Geometry)	16	15
6.1.4 (Point)	4	0
6.1.6 (Curve)	5	0
6.1.7 (LineString, Line, LinearRing)	2	0
6.1.8 (MultiCurve)	2	0
6.1.10 (Surface)	3	0
6.1.11 (Polygon, Triangle)	3	0
6.1.12 (PolyhedralSurface)	4	0
6.1.13 (MultiSurface)	3	0
6.1.15 (Relational Operators)	0	9
Total	42	24

Table 5.1: OGC Methods by Number of Required Geometries

all records and the new one

2. Read an existing record (by id)

PostGIS: select ... where id = Hadoop: map: emit only the record with the correct id; reduce: identity ClusterGIS: emit only the record with the correct id;

3. Update

PostGIS: update.... Hadoop: map: emit all records, updating the one with the correct id; reduce: identity ClusterGIS: emit all records updating the correct one

4. Destroy

PostGIS: delete ... Hadoop: map: emit all records except for the one with the correct id; reduce: identity ClusterGIS: emit all record except the correct one

5. Filter

PostGIS: select...where Hadoop: map: emit only records matching the requested criteria; reduce: identity ClusterGIS: emit only records matching the requested criteria

6. Find Nearest

PostGIS: Crazy set of queries (nested would have been nice, but too slow) Hadoop: map: loop through secondary dataset emitting id of primary and secondary; reduce: iden-

tity ClusterGIS: loop through primary and secondary dataset, emitting the id of primary and secondary which match

7. Chaining (Filter the datasets, then find the nearest in the filtered data)

PostGIS: Modify crazy set of queries to use the filter Hadoop: filter secondary dataset on load; map: only search for nearest for filtered data, skip others; reduce: identity ClusterGIS: remove nodes in the datasets not matching the filter, find nearest as before

5.2 Dataset Description

Full datasets; 34k employers, 1.2m parcels

sub datasets, how to generate from full

listing of datsets used

5.3 Hardware setup

Saguaro

5.4 implementations

5.4.1 PostGIS

5.4.2 Hadoop Prototype

5.4.3 ClusterGIS

CHAPTER 6   RESULTS

6.1   Performance Analysis

vary procs, maintain data

    vary data, maintain procs

    vary data, vary procs

6.2   Comparisons

PostGIS as baseline

    PostGIS to Hadoop

    Hadoop to ClusterGIS

    PostGIS to ClusterGIS

CHAPTER 7   RECOMMENDATIONS

Synthesis

7.1   Applicable Problem Spaces

Analysis of problem types that are good/bad for this approach

CHAPTER 8   CONCLUSION

By simplifying the requirements to handle process GIS operations in parallel, the cost to implementing parallel solutions will decrease, enabling more parallel processing methods to be created. This parallel methods will be able to take advantage of the increased processing powers made available through compute clusters.

Another benefit of this research is a classification of geospatial operations based on data access requirements and a sample set of queries to evaluate the efficiency of a parallel GIS processing implementation.

8.1   Future Work

additional decomposition methods, combined with alternative MPI communicators

addition of preprocessing methods (indexing, etc)

chunking of replicated dataset

CHAPTER Bibliography

[1] T. Ormsby and R. Burke, *Getting to Know ArcGIS Desktop: Basics of ArcView, ArcEditor, and ArcInfo*. ESRI, Inc., 2nd ed., 2004.

[2] J. Gray, "Quantum gis: the open-source geographic information system," *Linux J.*, vol. 2008, no. 172, p. 8, 2008.

[3] M. Neteler and H. Mitasova, *Open source GIS: a grass GIS approach*. Springer, 3rd ed., 2002.

[4] P. Waddell, A. Borning, M. Noth, N. Freier, M. Becke, and G. Ulfarsson, "Microsimulation of urban development and location choices: Design and implementation of urbansim," *Networks and Spatial Economics*, vol. 3, pp. 43–67, 2003.

[5] D. Blasby, "Building a spatial database in postgresql," in *Open Source Database Summit*, 2001.

[6] R. West, *Understanding ArcSDE*. ESRI, Inc., 2001.

[7] S. Ravada and J. Sharma, "Oracle8i spatial: Experiences with extensible databases," in *SSD '99: Proceedings of the 6th International Symposium on Advances in Spatial Databases*, (London, UK), pp. 355–359, Springer-Verlag, 1999.

[8] R. Pahle and N. Kerr, "A datacentric framework for research in planning," in *UPE8, The 8th International Symposium (UPE 8) of the International Urban Planning and Environment Association*, 2009.

[9] S. Guhathakurta, Y. Kobayashi, M. Patel, J. Holston, T. Lant, J. Crittenden, K. Li, G. Konjevod, and K. Date, "Digital phoenix project: A multidimensional journey through time." Not published, 2006.

[10] M. Haklay and P. Weber, "Openstreetmap: User-generated street maps," *Pervasive Computing, IEEE*, vol. 7, pp. 12–18, Oct.-Dec. 2008.

[11] D. DeWitt and J. Gray, "Parallel database systems: the future of high performance database systems," *Commun. ACM*, vol. 35, no. 6, pp. 85–98, 1992.

[12] J. Patel, J. Yu, N. Kabra, K. Tufte, B. Nag, J. Burger, N. Hall, K. Ramasamy, R. Lueder, C. Ellmann, J. Kupsch, S. Guo, J. Larson, D. De Witt, and J. Naughton, "Building a scaleable geo-spatial dbms: technology, implementation, and evaluation," in *SIGMOD '97: Proceedings of the 1997 ACM SIGMOD international conference on Management of data*, (New York, NY, USA), pp. 336–347, ACM, 1997.

[13] D. J. DeWitt, N. Kabra, J. Luo, J. M. Patel, and J.-B. Yu, "Client-server paradise," in *VLDB '94: Proceedings of the 20th International Conference on Very Large Data Bases*, (San Francisco, CA, USA), pp. 558–569, Morgan Kaufmann Publishers Inc., 1994.

[14] F. Huang, D. Liu, P. Liu, S. Wang, Y. Zeng, G. Li, W. Yu, J. Wang, L. Zhao, and L. Pang, "Research on cluster-based parallel gis with the example of parallelization on grass gis," in *Grid and Cooperative Computing, 2007. GCC 2007. Sixth International Conference on*, pp. 642–649, Aug. 2007.

[15] P. J. Braam, "Lustre file system," white paper, Cluster File Systems, Inc., 2007.

[16] G. Mackey, S. Sehrish, J. Bent, J. Lopez, S. Habib, and J. Wang, "Introducing map-reduce to high end computing," in *Petascale Data Storage Workshop, 2008. PDSW '08. 3rd*, pp. 1–6, Nov. 2008.

[17] "Opengis implementation specification for geographic information - simple feature access - part 1: Common architecture," standard, Open Geospatial Consortium, Inc., 2006.

[18] "Dbc/1012 data base computer concepts & facilities," tech. rep., Teradata Corp Document No C02-0001-00, 1983.

[19] J. Dean and S. Ghemawat, "Mapreduce: simplified data processing on large clusters," *Commun. ACM*, vol. 51, no. 1, pp. 107–113, 2008.